

TANMOY SAHA

Amarpur, Tripura, India

☎ +91 9366908415 ✉ contact@tanmoytssaha.cc [in](#) [Linkedin](#) [G](#) [Github](#) [globe](#) tanmoytssaha.cc

Software Engineer with 3+ years building data infrastructure and AI tooling at scale. Experienced in Python, Go, Docker, and AWS, with hands-on exposure to LLM-based agent systems and AI model deployment on platforms including RunPod. Strong communicator comfortable translating technical complexity for non-technical stakeholders.

TECHNICAL SKILLS

Languages: Python, SQL, Go, Scala, C++

AI and ML: MCP, LLM Orchestration, SDXL, Agentic Systems, RunPod

Distributed Systems and Data: Apache Spark, PySpark, Hive, Hadoop, Kafka

Cloud and Platform: AWS (S3, EMR, Glue, Lambda, Athena, EC2), Docker, Linux/Ubuntu

Data Architecture: ETL/ELT, Config-driven Pipelines, Data Modeling

Databases: PostgreSQL, SQL Server, MongoDB

Orchestration and DevOps: Step Functions, Airflow, Azkaban, Git

Observability: CloudWatch, Logging/Metrics/Alerting

Other: REST APIs, Django, DSA

EXPERIENCE

Paytm

Feb 2025 – Present

Data Engineer – Central Data Platform

Noida, Uttar Pradesh, India

- Reduced S3 storage by 45 percent by implementing and optimizing Life Cycle Policies (LCP), including removal of obsolete data and non-functional policies through strategic auditing and analysis of usage patterns.
- Boosted Spark job performance by narrowing data scans and eliminating redundant transformations, resulting in faster execution and reduced compute costs.
- Built an AI agent tool using the Model Context Protocol (MCP), integrated with Cursor AI and LLM-based orchestration, automating data pipeline deployment — cutting deployment time from 2 hours to 30 minutes.
- Refactored the batch data pipeline codebase by removing flat-file dependencies and centralizing metadata storage in a database, making it accessible through APIs for easier integration and maintenance.
- Improved internal monitoring dashboards by redesigning the UI and integrating additional metrics and alerting capabilities for better observability.
- Authored end-to-end technical documentation for the AI MCP Tool — covering architecture, integration guide, and usage workflows — and presented the solution to the engineering team and Director, driving org-wide adoption.

SpiceJet

Jan 2023 – Feb 2025

Data Engineer

Bangaluru, Karnataka, India

- Designed and implemented an end-to-end near real-time data pipeline on AWS using PySpark, EMR, Glue, and Lambda to support near real-time analytics.
- Worked on optimizing and adding new requirements on the existing pipeline.
- Collaborated with stakeholders to develop data visualization dashboards using Django, HTML, TailwindCSS, and JavaScript, providing actionable insights for decision-making.
- Documented the near real-time AWS data pipeline architecture (PySpark, EMR, Glue, Lambda) with data flow diagrams and runbooks, and delivered a technical presentation to cross-functional stakeholders and the Director to align teams on pipeline design and capabilities.

Wielabs

June 2022 – Jan 2023

Software Development Engineer Intern

Hyderabad, Telangana, India

- Built scalable backend services using Python, Django, and PostgreSQL with optimized query patterns and modular API design.
- Developed reusable UI components using Flutter, Tailwind, and JS to power a unified web/mobile experience.

PROJECTS

PantryAI – AI-Powered Kitchen App | *Personal Project*

- Built SDXL-based food image generation pipeline on RunPod serverless infrastructure for real-time inference.
- Architected a Go backend with Flutter frontend; managed full deployment lifecycle on cloud GPU infrastructure.

PRECOG – Flight DFDR Data Pipeline | *SpiceJet*

- Worked on scalable, multi-stage pipelines with PySpark to process flight DFDR data for analytics teams.
- Implemented data governance and quality checks to ensure clean, actionable insights for operations.

PRECOG – Flight DFDR Data Pipeline | *SpiceJet*

- Worked on scalable, multi-stage pipelines with PySpark to process flight DFDR data for analytics teams.
- Implemented data governance and quality checks to ensure clean, actionable insights for operations.

RateGain – Email Data Automation and Competitive Analysis | *SpiceJet*

- Automated extraction of CSV data from emails, saved as Parquet with partitioning for efficient querying.
- Leveraged Spark for preprocessing and built staging tables for cost and competition analysis.

EDUCATION

Kalinga University, Raipur

Bachelor of Technology in Computer Science (CPGA - 8.61)

Raipur, Chhattisgarh